

MiniRAG Assistant

Este proyecto permite hacer RAG (Retrieval-Augmented Generation) usando Chroma en Docker, embeddings de HuggingFace y consultas a un LLM de OpenAI.

IMPORTANTE: Antes de ejecutar el proyecto, asegúrate de tener **Docker Desktop instalado y abierto**, y que el servicio de Chroma esté corriendo (`docker compose up -d`).

Tabla de Contenidos

1. [Preparar el proyecto](#)
2. [Docker: Chroma remoto](#)
3. [Entorno Python](#)
4. [Variables de entorno](#)
5. [Coloca tus documentos](#)
6. [Indexador \(Chroma remoto + HF embeddings\)](#)
7. [Inspección de colección](#)
8. [App de Streamlit \(RAG con LLM\)](#)
9. [Resetear contenido](#)
10. [Problemas comunes & soluciones](#)

1) Preparar el proyecto

Estructura de proyecto

```
mkdir mini-rag-assistant `
; cd mini-rag-assistant `
; mkdir docs `
; New-Item app.py -ItemType File `
; New-Item indexer.py -ItemType File `
; New-Item verify_chroma.py -ItemType File `
; New-Item inspect_chroma.py -ItemType File `
; New-Item reset_chroma.py -ItemType File `
; New-Item requirements.txt -ItemType File `
; New-Item .env -ItemType File `
; New-Item docker-compose.yml -ItemType File
```

2) Docker: Chroma remoto

```
services:
  chroma:
    image: chromadb/chroma:latest
```

```
container_name: chroma
environment:
  CHROMA_SERVER_HOST: 0.0.0.0
  CHROMA_SERVER_HTTP_PORT: 8000
  ALLOW_RESET: "true"
ports:
  - "8000:8000"
volumes:
  - ./chroma_data:/chroma/.chroma
restart: unless-stopped
```

```
docker compose up -d
```

3) Entorno Python

```
py -3.11 -m venv venv
.\venv\Scripts\Activate
```

```
chromadb==1.0.20
langchain==0.3.27
langchain-core==0.3.74
langchain-community==0.3.27
langchain-chroma==0.2.5
langchain-huggingface==0.3.1
langchain-openai==0.3.32
langchain-text-splitters==0.3.9
streamlit==1.37.1
python-dotenv==1.0.1
pypdf==4.2.0
sentence-transformers==2.6.1
```

```
pip install -r requirements.txt
```

4) Variables de entorno

```
OPENAI_API_KEY=tu_api_key_sin_comillas
```

5) Coloca tus documentos

Crea TXT de prueba en la carpeta `docs/`.

6) Indexador (Chroma remoto + HF embeddings)

```
# indexer.py completo
" + `import os
import time
import chromadb
from chromadb.config import Settings
from dotenv import load_dotenv
from langchain_huggingface import HuggingFaceEmbeddings
from langchain_text_splitters import RecursiveCharacterTextSplitter
from langchain_community.document_loaders import TextLoader, PyPDFLoader

load_dotenv()

HOST = "localhost"
PORT = 8000
COLLECTION = "mini_rag"
DOCS_DIR = "docs"

def load_docs(docs_dir: str):
    if not os.path.exists(docs_dir):
        print(f"❌ No existe la carpeta '{docs_dir}'.")
        return []

    files = [f for f in os.listdir(docs_dir) if f.lower().endswith((".txt",
".pdf"))]
    if not files:
        print(f"❌ No hay .txt/.pdf en '{docs_dir}'.")
        return []

    print(f"💣 Archivos en '{docs_dir}': {files}")
    docs = []
    for fname in files:
        fpath = os.path.join(docs_dir, fname)
        if fname.lower().endswith(".txt"):
            loader = TextLoader(fpath, encoding="utf-8")
        else:
            loader = PyPDFLoader(fpath)
        loaded = loader.load()
        print(f"🌙 '{fname}' → {len(loaded)} documento(s)")
        docs.extend(loaded)
    return docs

def main():
```

```

documents = load_docs(DOCS_DIR)
if not documents:
    return

print(f"1 📄 Dividiendo {len(documents)} documentos en chunks...")
splitter = RecursiveCharacterTextSplitter(chunk_size=600,
chunk_overlap=80)
chunks = splitter.split_documents(documents)
print(f"✅ Total de chunks: {len(chunks)}")
if not chunks:
    print("❌ No se generaron chunks.")
    return

print("9 📦 Cargando modelo de embeddings (HF MiniLM-L6-v2)...")
embedder = HuggingFaceEmbeddings(model_name="sentence-transformers/all-
MiniLM-L6-v2")
test_vec = embedder.embed_query("prueba")
print(f"📏 Dimensión de embedding (query): {len(test_vec)}")

print("🌐 Conectando a Chroma remoto...")
client = chromadb.HttpClient(host=HOST, port=PORT,
settings=Settings(anonymized_telemetry=False))
print(f"✅ Conectado")

print(f"☢️ Preparando colección '{COLLECTION}'...")
try:
    client.delete_collection(COLLECTION)
    print("6 🗑️ Colección previa eliminada")
except Exception:
    print(f"📄 Colección previa no existía, continuamos")

collection = client.get_or_create_collection(name=COLLECTION)
print(f"✅ Colección lista")

print(f"📁 Insertando chunks en batch...")
BATCH = 64
ids, texts = [], []
for i, c in enumerate(chunks, 1):
    ids.append(f"doc_{i}")
    texts.append(c.page_content)

    if len(ids) == BATCH or i == len(chunks):
        print(f"➡️ Batch con {len(ids)} items... (generando embeddings)")
        embs = embedder.embed_documents(texts)
        collection.add(ids=ids, documents=texts, embeddings=embs)
        print(f"✅ Upsert de {len(ids)} items")
        ids, texts = [], []

time.sleep(0.4)
count = collection.count()
print(f"!!! Indexación completa. Documentos en '{COLLECTION}': {count}")

```

```
if __name__ == "__main__":  
    main()
```

python indexer.py

7) Inspección de colección

```
# inspect_chroma.py completo  
import chromadb  
from chromadb.config import Settings  
  
client = chromadb.HttpClient(host="localhost", port=8000,  
settings=Settings(anonymized_telemetry=False))  
COLLECTION_NAME = "mini_rag"  
  
try:  
    collection = client.get_collection(COLLECTION_NAME)  
except Exception as e:  
    print(f"❌ No se encontró la colección '{COLLECTION_NAME}'. Error: {e}")  
    exit()  
  
print(f"🔍 Inspeccionando colección: {COLLECTION_NAME}")  
count = collection.count()  
print(f"📊 Total de documentos: {count}")  
  
if count == 0:  
    print(f"⚠️ La colección está vacía.")  
    exit()  
  
batch_size = 5  
for offset in range(0, count, batch_size):  
    results = collection.get(include=["documents", "metadatas"],  
limit=batch_size, offset=offset)  
    for i, doc in enumerate(results["documents"]):  
        doc_id = results["ids"][i]  
        meta = results["metadatas"][i]  
        preview = doc[:100].replace("\n", " ") + "..." if len(doc) > 100 else doc  
  
        print(f"\n📄 ID {doc_id}")  
        print(f"📄 Texto: {preview}")  
        print(f"📄 Metadatos: {meta}")
```

python inspect_chroma.py

8) App de Streamlit (RAG con LLM)

```
# app.py completo
import streamlit as st
import chromadb
from chromadb.config import Settings
from dotenv import load_dotenv
from langchain_huggingface import HuggingFaceEmbeddings
from langchain_chroma import Chroma
from langchain_openai import ChatOpenAI
from langchain.chains import RetrievalQA

load_dotenv()
HOST = "localhost"
PORT = 8000
COLLECTION = "mini_rag"

st.set_page_config(page_title="Mini RAG Assistant", page_icon="📖")
st.title("📖 Mini RAG Assistant")
st.write("Haz preguntas sobre tus documentos indexados en Chroma remoto.")

query = st.text_input("Tu pregunta:")

if query:
    with st.spinner("Procesando..."):
        st.write("♦ Inicializando embeddings (HF MiniLM-L6-v2)...")
        embeddings = HuggingFaceEmbeddings(model_name="sentence-transformers/
all-MiniLM-L6-v2")
        st.success("✅ Embeddings inicializados")
        st.write("♦ Conectando a Chroma remoto...")
        client = chromadb.HttpClient(host=HOST, port=PORT,
settings=Settings(anonymized_telemetry=False))
        st.success("✅ Conexión con Chroma exitosa")
        st.write("♦ Cargando vectorstore...")
        vectorstore = Chroma(client=client, collection_name=COLLECTION,
embedding_function=embeddings)
        retriever = vectorstore.as_retriever(search_kwargs={"k": 3})
        st.success("✅ Retriever listo")
        st.write("♦ Inicializando LLM...")
        llm = ChatOpenAI(model="gpt-4o-mini", temperature=0)
        st.success("✅ LLM listo")
        st.write("♦ Construyendo chain de RAG...")
        qa_chain = RetrievalQA.from_chain_type(llm=llm, retriever=retriever,
return_source_documents=True)
        st.success("✅ Chain creada")
        st.write("♦ Consultando el modelo...")
        result = qa_chain.invoke({"query": query})
        st.success("✅ Consulta completada")
        st.markdown("### Respuesta")
        st.write(result["result"])
```

```

st.markdown("### 🕒 Documentos fuente")
src_docs = result.get("source_documents", [])
if src_docs:
    for i, doc in enumerate(src_docs, 1):
        st.markdown(f"***Documento {i}:***")
        st.write(doc.page_content[:600] + "...")
        st.caption(f"Metadata: {doc.metadata}")
else:
    st.warning("⚠️ No se encontraron documentos relevantes.")
st.markdown("### 🕒 Top-K por similitud (debug)")
sims = vectorstore.similarity_search_with_score(query, k=3)
for j, (d, score) in enumerate(sims, 1):
    st.write(f"***{j}.** score={score:.4f}")
    st.caption(d.page_content[:200] + "...")

```

```
streamlit run app.py
```

9) Resetear contenido

```

# reset_chroma.py completo
import chromadb
from chromadb.config import Settings
HOST = "localhost"
PORT = 8000
COLLECTION = "mini_rag"
client = chromadb.HttpClient(host=HOST, port=PORT,
settings=Settings(anonymized_telemetry=False))
try:
    client.delete_collection(COLLECTION)
    print(f"🗑️ Colección '{COLLECTION}' eliminada.")
except Exception as e:
    print("📄 La colección no existía:", e)
# client.reset() # opción B, opcional

```

```
python reset_chroma.py
```

10) Problemas comunes & cómo solucionarlos

- Dimensión de embeddings no coincide (384 vs 1536)
- Streamlit "se desconecta" o cuelga
- ImportError: cannot import name 'Chroma'... tras crear un inspect.py
- No logra borrar carpeta local por Docker